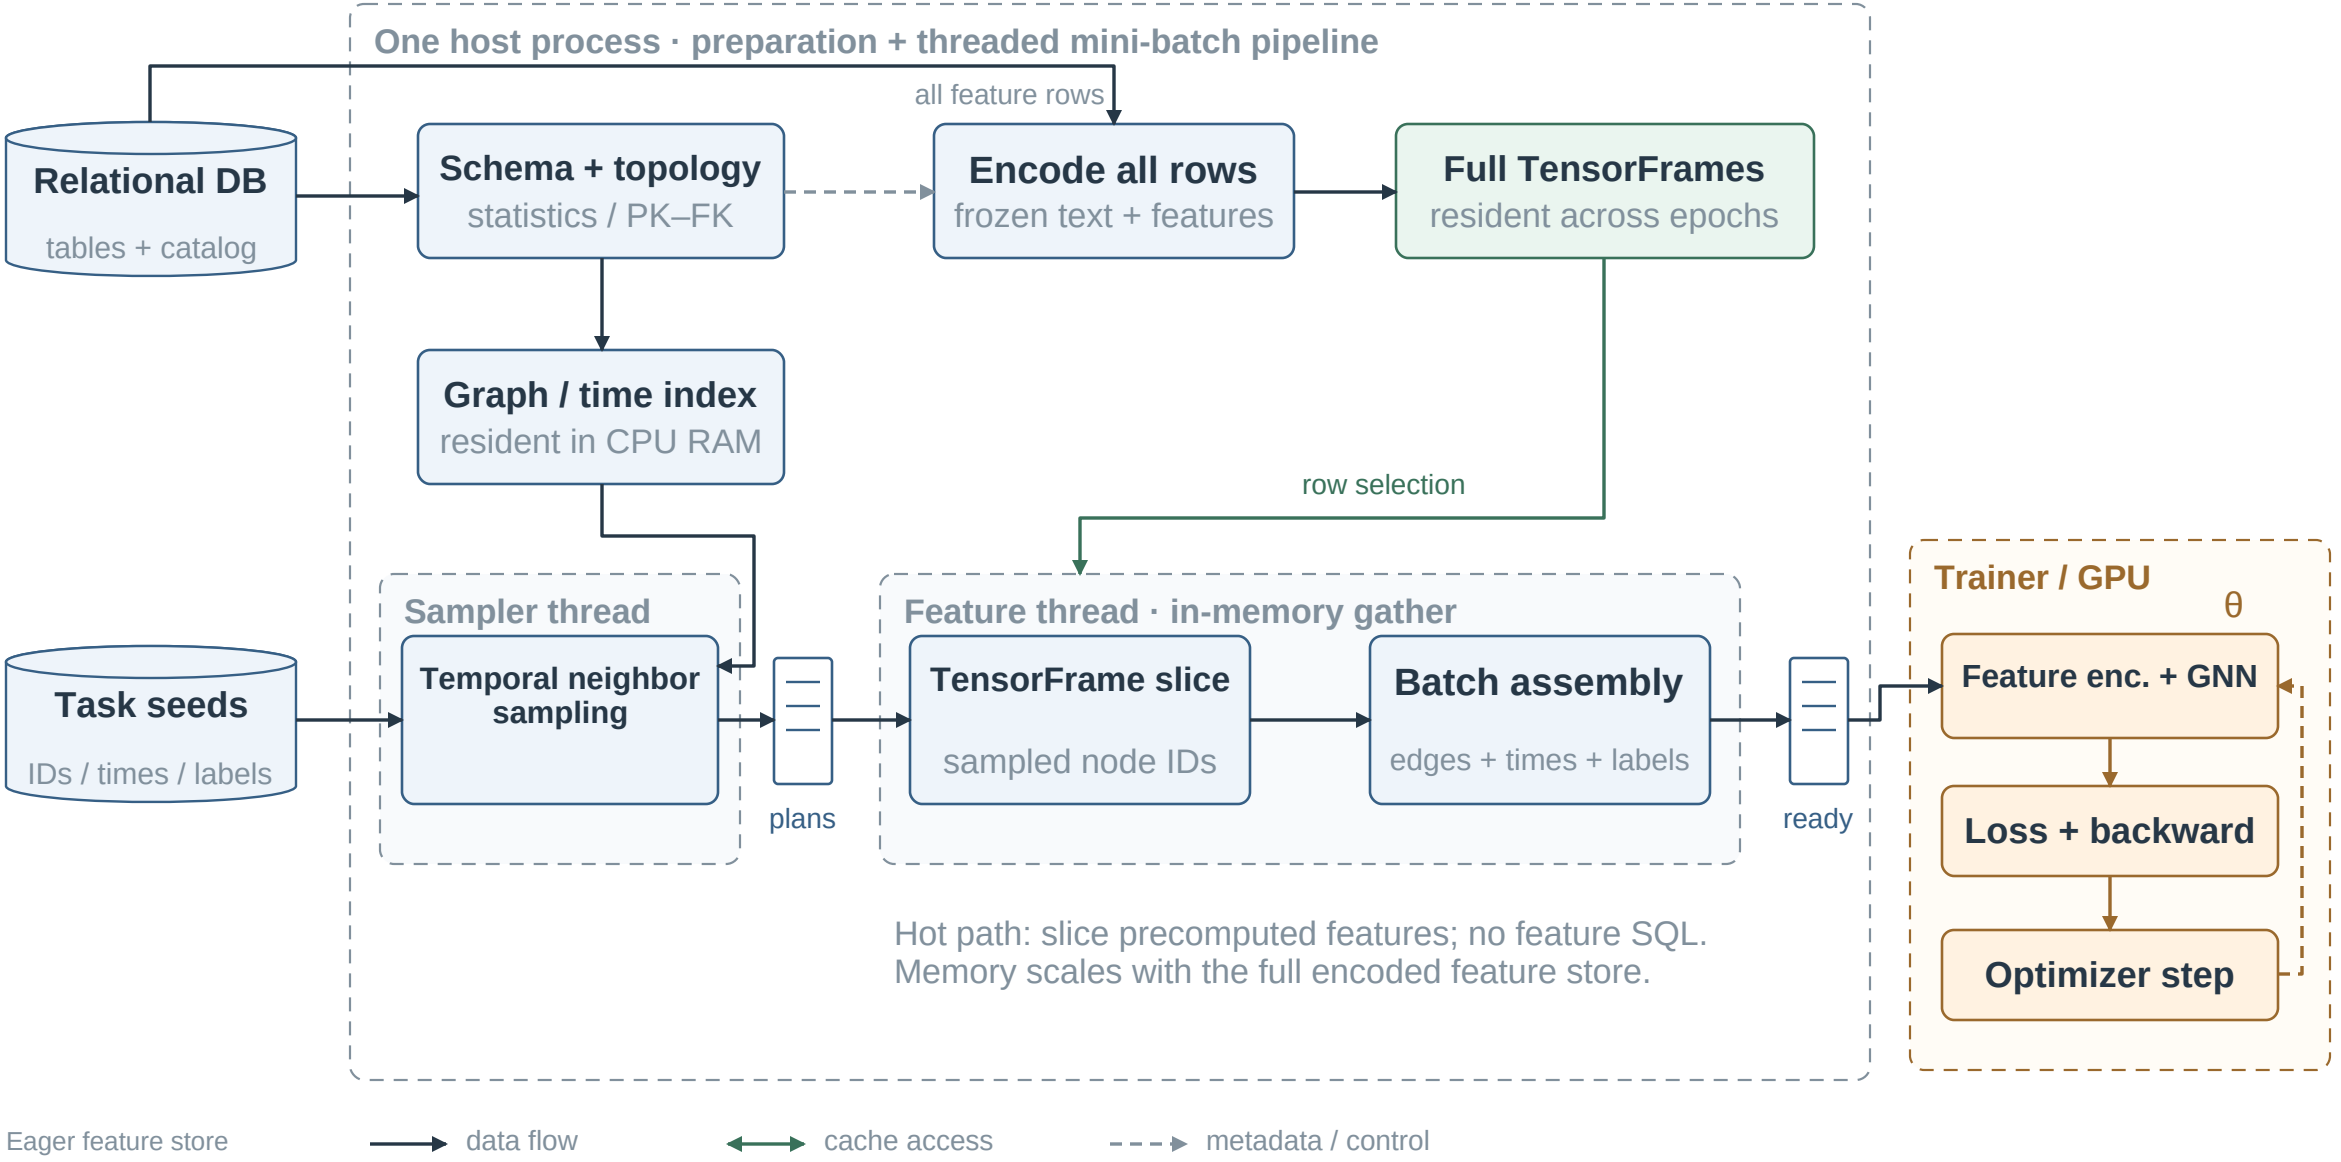
One host process · preparation + threaded mini-batch pipeline
Relational DB
tables + catalog
Schema + topology
statistics / PK–FK
all feature rows
Encode all rows
frozen text + features
Full TensorFrames
resident across epochs
Graph / time index
resident in CPU RAM
row selection
Sampler thread
Task seeds
IDs / times / labels
Temporal neighbor sampling
plans
Feature thread · in-memory gather
TensorFrame slice
sampled node IDs
Batch assembly
edges + times + labels
ready
Trainer / GPU
θ
Feature enc. + GNN
Loss + backward
Optimizer step
Hot path: slice precomputed features; no feature SQL.
Memory scales with the full encoded feature store.
Eager feature store
data flow
cache access
metadata / control